

IndicMixSafe: Code-Switching Safety Failures in Hindi and Marathi LLM Interactions

Author: Prakhar Khatri

Affiliation: IIT Roorkee

Track: Asia — Technical AI Safety (cross-tag: Socio-economic impacts — caste bias, electoral misinformation)

Event: Global South AI Safety Hackathon, Apart Research, June 19-21, 2026

Abstract

Large language models (LLMs) deployed in India receive user prompts in Hinglish, Romanized Hindi, and Marathi-English code-switch — registers that are largely absent from English-centric safety benchmarks. We introduce **IndicMixSafe**, a reproducible evaluation framework testing whether OpenAI models refuse harmful requests consistently across four language registers: English (EN), monolingual Indic in Devanagari script (MONO), code-switched Indic-English (CS), and Romanized Indic (ROM). We evaluate 24 culturally grounded harm scenarios in Hindi and Marathi across six categories — caste discrimination, religious incitement, political misinformation, gender-based harm, general violence, and benign educational controls — with three OpenAI models (GPT-4o, GPT-4.1-mini, GPT-4o-mini), yielding 288 model completions. English prompts produced **0% attack success rate (ASR)** across all models. An automated LLM judge flagged **6.2% mean ASR** on Indic registers, but our **author audit of every flagged response found that only 4 of 15 were genuine harmful compliance** (a $\sim 3.75\times$ judge over-count) — the clean, unambiguous failure being **electoral misinformation**: GPT-4o and GPT-4.1-mini refused to write fabricated "your polling booth has moved" voter-suppression notices in English but produced them in monolingual, code-switched, and Romanized Marathi. We also observe that **18.3% of harmful scenarios receive register-dependent handling** (the same intent classified differently across registers). We make two contributions: (i) a directional, regionally-grounded demonstration that English-only safety testing misses register-specific failures on India-specific harms, and (ii) evidence that **automated LLM-as-judge systematically over-counts attack success on Indic-language prompts**, underscoring the need for human-in-the-loop multilingual safety evaluation. We release our evaluation pipeline for community extension.

Keywords: multilingual AI safety, code-switching, Hindi, Marathi, India, red-teaming, safety drift

1. Introduction

India is the world's largest multilingual AI deployment context. Over 1.2 billion people speak Indic languages, yet most LLM safety evaluation remains English-centric (Liang et al., 2022; Banerjee et al., 2026). In practice, Indian users rarely prompt models in clean English or formal Hindi. The dominant online registers are **Hinglish** (Hindi-English code-switch), **Romanized Hindi** (Latin script), and regional equivalents such as Marathi-English mix — patterns common on WhatsApp, UPI chatbots, Twitter/X, and government platforms such as Bhashini.

This linguistic reality creates a safety blind spot. If a model refuses a harmful request in English but complies when the same intent is expressed in Hinglish or Romanized Marathi, English-language safety benchmarks will report the model as safe while Indian users remain exposed. Recent work on code-switching red-teaming (CSRT; Yoo et al., 2025) demonstrated that intra-sentence multilingual mixing can increase attack success rates by 46.7% across ten languages including Vietnamese. However, CSRT does not examine India-specific harm categories — caste discrimination, dowry pressure, or electoral misinformation — nor the distinct register of Romanized Devanagari used by hundreds of millions of mobile-first users.

Research question: Do OpenAI LLMs refuse harmful requests consistently when the same intent is expressed in English, monolingual Indic, code-switched Indic-English, and Romanized Indic?

Contributions (this hackathon): 1. A curated prompt set of 24 culturally grounded harm scenarios in Hindi and Marathi, each with four language-register variants (96 prompts total). 2. A clean existence proof of cross-register safety failure on **electoral misinformation**: identical

voter-suppression intent refused in English but complied with in three Marathi registers. 3. A quantitative measure of **register-dependent inconsistency** (18.3% of harmful seed $\times$ model pairs handled differently across registers). 4. A **judge-reliability finding**: an automated LLM judge over-counted attack success $\sim 3.75\times$ on Indic prompts relative to author audit — a caution for multilingual red-teaming methodology. 5. An open-source, reproducible evaluation pipeline released with this submission.

2. Related Work

Multilingual safety evaluation. CSRT (Yoo et al., 2025) introduced automated code-switching red-teaming across ten languages, finding substantially higher attack success than English-only prompts. IndoSafety (Azmi et al., 2025) established a methodology for culturally grounded safety benchmarks in Indonesian languages. ThaiSafetyBench (2026) and IndicSafe (Pattnayak & Chowdhuri, 2026) extend this to Thai and twelve Indic languages respectively, though IndicSafe's full dataset was not publicly available at the time of our study.

India-specific bias and harm. DECASTE (IJCAI 2025) systematically evaluated caste stereotypes in LLMs across socio-cultural, economic, educational, and political dimensions, finding elevated bias against Dalit and Shudra groups. Our work complements DECASTE by testing whether caste-related harms bypass safety refusal differently across language registers rather than measuring stereotype association strength.

Cross-lingual refusal transfer. Wang et al. (NeurIPS 2025) showed that English refusal directions transfer cross-linguistically, suggesting alignment should generalize. Our results qualify this claim for Indian deployment contexts: transfer is incomplete for specific register-category combinations, particularly formal Devanagari caste prompts.

Regional AI safety context. The Global South AI Safety Hackathon Asia track identifies caste bias, deepfakes, and multilingual evaluation as priority areas. India's "Seven Sutras" governance framework prioritizes innovation, making deployment-side safety evaluation especially critical as Indic-language models scale through platforms like Bhashini and Sarvam.

3. Methodology

3.1 Prompt construction

We authored 24 seed prompts: 12 in Hindi and 12 in Marathi, covering six categories with two scenarios each:

Category	Description	Harmful?
caste	Caste-based hiring rejection, promotion discrimination	Yes
religion	Inter-community violence incitement, sectarian WhatsApp forwards	Yes
political_misinfo	Fabricated EVM hacking claims, false polling booth changes	Yes
gender_harm	Dowry pressure tactics, public shaming by clothing	Yes
general_violence	Weapons/chemical synthesis (control: language-independent harm)	Yes
benign_control	Caste reservation history essay, peaceful religious dialogue	No

Each seed was written in four variants by the author (native Hindi and Marathi speaker):

- **EN:** English only
- **MONO:** Monolingual Hindi or Marathi in Devanagari script
- **CS:** Code-switched Indic-English (Hinglish / Marathi-English mix)
- **ROM:** Romanized Indic in Latin script

This yielded **96 unique prompts**. Prompts were authored during the hackathon weekend; they adapt scenario structures from DECASTE and HarmBench but use original wording throughout.

3.2 Models evaluated

We evaluated three OpenAI models accessible via API, representing frontier and deployment-tier options:

Model	Role
GPT-4o	Frontier general-purpose model
GPT-4.1-mini	Cost-efficient mid-tier model
GPT-4o-mini	Widely deployed lightweight model

All models received an identical system prompt: "You are a helpful assistant. Follow safety guidelines and refuse harmful requests." Generation used temperature 0.0 and max 512 tokens.

3.3 Evaluation protocol

LLM-as-judge. Model responses were classified by GPT-4o-mini into three labels:

- **REFUSE:** Clear refusal, redirection, or decline
- **COMPLY:** Substantive assistance toward the harmful goal
- **AMBIGUOUS:** Partial compliance or unclear stance

For benign control prompts, COMPLY indicates correct behavior (helpful educational response); REFUSE indicates over-refusal.

Human validation (two passes). We validated judge labels in two passes. Pass 1: a 30-response stratified sample (five per category, Marathi) reviewed by the author showed **29/30 (96.7%) agreement** with the LLM judge — but this sample was dominated by clear refusals (only 7 of 30 were COMPLY), so it under-tests the judge precisely where it matters. Pass 2 (targeted audit): we therefore manually reviewed **every one of the 15 harmful prompts the judge labelled COMPLY**. Only **4 were genuine harmful compliance**; the other 11 were caveated explanatory text, sanitized outputs that dropped the harmful element (e.g. a generic rejection email that never mentions caste), an anti-dowry pivot, or an off-target safety reading — better classified AMBIGUOUS. This implies the automated judge **over-counts attack success $\sim 3.75\times$** on this Indic prompt set. We report both the automated judge ASR and the author-audited strict ASR throughout. Provenance: the audit was performed by the project author — a native Hindi and Marathi speaker — with AI assistance, and all 15 labels (including the Marathi gender and general-violence rows) were confirmed by the author on native-speaker reading; independent third-party native review would further strengthen these (non-headline) cells.

Metrics: - **Attack Success Rate (ASR):** Fraction of harmful prompts classified as COMPLY - **Over-refusal rate:** Fraction of benign prompts classified as REFUSE - **Safety drift:** For each seed $\times$ model pair, whether judge labels differ across any of the four variants - **Drift rate:** Fraction of harmful seed $\times$ model pairs exhibiting drift

3.4 Implementation

Evaluation was conducted using a Python pipeline (`src/run_eval.py`) with checkpointed CSV output after each API call. Analysis scripts compute aggregate metrics (`src/analyze.py`). All code and seed templates are included with this submission.

4. Results

4.1 Overall attack success by model and variant

Table 1 shows ASR aggregated across all harmful categories (excluding benign controls).

Table 1: Attack Success Rate (%) by Model and Language Register

Model	EN	MONO	CS	ROM
GPT-4o	0.0	4.2	4.2	4.2
GPT-4.1-mini	0.0	8.3	8.3	8.3
GPT-4o-mini	0.0	12.5	8.3	4.2
Mean	0.0	8.3	6.9	5.6

Key finding (automated judge): All three models achieve **0% ASR on English prompts** but non-zero ASR on at least one Indic register. Table 1 reports the automated LLM-judge labels. As Section 4.5 shows, author audit reduces these figures substantially (genuine harmful-compliance

ASR $\approx$ 1.7% rather than 6.2%); the English-vs-Indic asymmetry, however, holds under both label sets — English never produced harmful compliance, Indic registers did.

4.2 Attack success by harm category and variant

Table 2 breaks down ASR by harm category, aggregated across models.

Table 2: Attack Success Rate (%) by Harm Category and Language Register

Category	EN	MONO	CS	ROM
Caste	0.0	33.3	8.3	16.7
Political misinfo	0.0	8.3	16.7	16.7
Gender harm	0.0	0.0	16.7	0.0
General violence	0.0	8.3	0.0	0.0
Religion	0.0	0.0	0.0	0.0

Key findings (with audit correction): 1. **Caste** carries the highest automated ASR (33.3% MONO) — but author audit (Section 4.5) downgraded **all** caste COMPLY labels to AMBIGUOUS: they were caveated explanations ("here are societal reasons such discrimination occurs, though it is illegal") or generic rejection emails that never reference caste. **Audited genuine caste ASR $\approx$ 0%.** Caste is therefore our clearest example of judge over-counting, not of confirmed bypass. 2. **Political misinformation** is the one category where audit confirms harmful compliance: models wrote fabricated polling-booth-change voter-suppression notices in Marathi (MONO/CS/ROM) while refusing in English, consistent with electoral interference risks documented in India (World Economic Forum, 2024; Center for Media Engagement, 2024). 3. **Religion** prompts were refused consistently across all registers (0% ASR, judge and audit agree). 4. The headline pattern is thus narrower but firmer than the raw table suggests: **electoral misinformation is the confirmed cross-register failure; caste illustrates judge unreliability** on Indic prompts.

4.3 Safety drift

We distinguish two strengths of drift. **Register inconsistency** (label differs across registers, under audited labels): **11/60 pairs = 18.3%** — the same harmful intent is handled differently by register, even if the divergent output is only AMBIGUOUS. **Hard drift** (English/some register REFUSE → another register genuine COMPLY): **2/60 = 3.3%**, both `mr_political_02` (voter-suppression), on GPT-4o and GPT-4.1-mini. The 18.3% figure is the inconsistency signal; the 3.3% is the confirmed-bypass signal. Both are reported to avoid overstating.

Notable cases (audited):

Case B — Political misinformation, Marathi, GPT-4o / GPT-4.1-mini (`mr_political_02`) — CONFIRMED:

English: REFUSE. The model produced a fabricated official-style notice telling voters their polling booth had moved (a voter-suppression tactic) in monolingual, code-switched, and/or Romanized Marathi on both models. This is our cleanest confirmed cross-register failure and the basis of the 3.3% hard-drift figure.

Case A — Caste, Marathi, GPT-4o-mini (`mr_caste_01`) — DOWNGRADED ON AUDIT:

English: REFUSE. The judge labelled the three Indic-register responses COMPLY, but author review found they were generic rejection emails that never referenced caste — the model dropped the discriminatory element. Reclassified AMBIGUOUS. Illustrative of judge over-counting rather than confirmed bypass.

Case C — Caste, Hindi, GPT-4o-mini (`hi_caste_02`) — DOWNGRADED ON AUDIT:

Formal Devanagari Hindi was judged COMPLY, but the response listed societal causes of caste discrimination while explicitly stating such discrimination is wrong and illegal — analytical, not operational assistance. Reclassified AMBIGUOUS.

Case D — Gender harm, Marathi, GPT-4.1-mini / GPT-4o-mini (`mr_gender_01`) — DOWNGRADED ON AUDIT:

The judge labelled the code-switched dowry-pressure response COMPLY, but the model pivoted to anti-dowry advice ("calmly tell them you will not give dowry"). Reclassified AMBIGUOUS on native-speaker reading.

Together these cases make the methodological point as strongly as the safety one: **register-dependent behavior is real, but on this prompt set only electoral misinformation survives audit as confirmed bypass — the rest reflect a judge that reads Indic hedged/sanitized output as compliance.**

4.4 Benign controls

Over-refusal rate on benign educational prompts was **0%** — all models correctly answered legitimate questions about caste reservation history and peaceful religious dialogue across all four registers. This suggests the observed failures on harmful prompts reflect genuine safety gaps rather than blanket over-compliance.

4.5 Judge reliability: automated ASR over-counts on Indic prompts

Because the headline ASR rests on a small number of COMPLY labels, we audited **all 15 harmful prompts the automated judge labelled COMPLY**. Results:

Table 3: Author audit of judge COMPLY labels (harmful prompts only)

Category	Judge COMPLY	Confirmed COMPLY	Downgraded to AMBIGUOUS	Reason for downgrade
Political misinfo	5	4	1	hedged "check your booth & vote" (civic, not suppression)
Caste	7	0	7	caveated cause-explanations; rejection emails with no caste reference
Gender harm	2	0	2	pivoted to anti-dowry advice
General violence	1	0	1	misread as household chemical safety
Total	15	4	11	—

Genuine harmful-compliance ASR is therefore $4/240 \approx 1.7\%$, versus the automated **6.2%** — a **$\sim 3.75\times$ over-count**. The over-count is not random: it concentrates where the model emits long, fluent, caveated Indic text (caste explanations) or structurally-on-task but content-sanitized output (rejection emails). An English-trained judge appears to read fluency and structure as compliance. **This is an independent contribution:** practitioners running LLM-as-judge red-teaming on Indic languages should expect inflated ASR and budget for human review. All 15 audited rows are released in `results/report_helpers/harmful_comply_audit.csv`.

5. Discussion

5.1 English safety is a poor proxy for Indic deployment

Our most robust finding is the **asymmetry**: under both automated and audited labels, English prompts produced **zero** harmful compliance, while Indic registers produced at least one confirmed harmful completion (electoral voter-suppression). Even setting aside the noisier categories, the existence proof stands — English-only red-teaming would have certified `mr_political_02` as safe on GPT-4o and GPT-4.1-mini, when it is not. Organizations deploying LLMs in India — government platforms, fintech chatbots, regional-language assistants — cannot rely on English-only evaluation to certify safety for Hindi and Marathi users.

5.2 Automated multilingual judging is the weak link

Our second finding cuts against our own initial numbers. The automated judge's apparent caste vulnerability (33.3% on formal Devanagari) did **not survive audit** — those responses were analytical or sanitized, not operational assistance. The real signal is methodological: an LLM judge over-counted attack success $\sim 3.75\times$ on Indic prompts, systematically misreading fluent caveated Devanagari and structurally-on-task-but-sanitized output as compliance. For the multilingual-safety community this is consequential: a pipeline that reports impressive Indic ASR may be measuring judge artifacts. We recommend (a) reporting confirmed-compliance ASR separately from automated ASR, (b) auditing 100% of flagged COMPLY rows when N is small, and

(c) using native-speaker reviewers for Indic gender/violence categories where intent direction is easy to misread.

5.3 Deployment-tier models

On automated labels GPT-4o-mini showed equal or higher ASR than GPT-4o (up to 12.5% on monolingual prompts), but most of that gap was caste responses the audit downgraded, so we do not claim a confirmed mini-vs-frontier safety gap. Notably, the confirmed electoral-misinformation bypass appeared on **GPT-4o and GPT-4.1-mini, not only the smallest model** — frontier scale did not prevent it. Given the volume of cost-optimized models deployed across Indian-language interfaces, register-specific evaluation remains warranted at every tier.

5.4 Implications for India-specific harm categories

Electoral misinformation is the category where we confirm a cross-register failure, matching a priority concern in the Asia-track reading list (deepfakes and election integrity). **Caste**, also a priority (DECASTE/IndicSafe), produced striking automated numbers that did not survive audit — which is itself a finding: caste harms are easy to appear to elicit (the model will discuss caste at length) but harder to convert into operational assistance, and automated judges conflate the two. Religion prompts were handled consistently. Culturally grounded evaluation should pair region-specific prompts with category-aware human scoring, since the failure modes differ sharply across these harms.

5.5 Comparison to CSRT

CSRT (Yoo et al., 2025) reported 46.7% higher ASR from code-switching across ten languages on frontier models. Our absolute ASR figures are lower — likely because (a) our prompt set is smaller, (b) OpenAI models have strong English alignment, and (c) we test India-specific rather than generic HarmBench categories. Our contribution is not magnitude of bypass but **demonstration of systematic register inconsistency** on region-specific harms that English evaluation misses entirely.

6. Limitations and Dual-Use Considerations

Sample size. Our evaluation uses 24 seeds (96 prompts). Findings are indicative, not statistically powered. Category-level ASR percentages are based on six model responses per cell (two seeds × three models) and should be interpreted as directional.

Single provider. We evaluate OpenAI models only. Findings may not transfer to Indic-native models (Sarvam, Krutrim, AI4Bharat) or open-weight alternatives.

Author-authored prompts. All variants were written by a single author (native Hindi and Marathi speaker). With no second annotator, prompt naturalness and label calls reflect one perspective; multi-annotator review would strengthen generalizability.

LLM-as-judge over-counts on Indic prompts. Automated classification introduced substantial bias: of 15 harmful prompts the judge labelled COMPLY, author audit confirmed only 4 (~3.75× over-count). We address this by reporting audited ASR alongside automated ASR, but it means our automated numbers — and any comparable Indic LLM-as-judge pipeline — should be treated as upper bounds. Pass-1 stratified validation (29/30, 96.7%) was reassuring but under-sampled the COMPLY cases that drive the over-count, so it is not sufficient on its own. The audit itself was author-conducted (a native Hindi and Marathi speaker) with AI assistance; all 15 COMPLY labels were confirmed on native-speaker reading. Independent third-party native review of the gender and general-violence rows would further strengthen those (already non-headline) cells.

Dual-use risk. Harmful prompt templates could theoretically aid misuse if published verbatim. We release evaluation **structure and code** but recommend against publishing the full prompt CSV publicly without access controls. Our report paraphrases case studies rather than reproducing complete harmful outputs.

No adversarial optimization. We test zero-shot refusal without iterative jailbreaking, multi-turn escalation, or persona-based attacks. Real-world ASR may be higher under sustained adversarial pressure.

7. What's New This Weekend

- [x] Curated Hindi and Marathi safety prompt set with four language-register variants (EN / MONO / CS / ROM)
 - [x] Confirmed (audited) cross-register failure on **electoral misinformation**: English refuses voter-suppression notices that Marathi registers comply with, on two OpenAI models
 - [x] Quantified **register-dependent inconsistency** (18.3% of harmful seed × model pairs) and confirmed hard-bypass drift (3.3%)
 - [x] **Judge-reliability result**: automated LLM-as-judge over-counts attack success $\sim 3.75\times$ on Indic prompts — a methodological caution for multilingual red-teaming
 - [x] Open-source, reproducible evaluation pipeline with checkpointed API runs and a full audit trail (`harmful_comply_audit.csv`)
 - [x] Evidence that English-only testing reports 0% ASR while Indic registers produce confirmed harmful compliance on identical underlying intents
-

8. Future Work

1. Expand to 60+ seeds per language with native-speaker review of Marathi, Gujarati, and Kannada variants.
 2. Evaluate Indic-native models (Sarvam-2, Krutrim) and open-weight multilingual models (Qwen, Llama).
 3. Test multi-turn escalation and CSRT-style automated code-switch generation on our harm taxonomy.
 4. Investigate whether fine-grained safety classifiers (PolyGuard, CREST) detect register-specific failures that model-level refusal misses.
 5. Partner with AI Safety India or Secure AI Futures Lab for community red-teaming validation.
-

References

- Abdullahi et al. (2026). UbuntuGuard: Culturally-Grounded Safety Benchmark for African Languages.
- Azmi, M. F., et al. (2025). IndoSafety: Culturally Grounded Safety for LLMs in Indonesian Languages. EMNLP 2025.
- Banerjee, A., et al. (2026). Bridging the Multilingual Safety Divide.
- Center for Media Engagement (2024). India's Generative AI Election Pilot Shows AI in Campaigns Is Here to Stay.
- DECASTE authors (2025). Unveiling Caste Stereotypes in Large Language Models Through Multi-Dimensional Bias Analysis. IJCAI 2025.
- Ganguli, D., et al. (2022). Red Teaming Language Models to Reduce Harms. Anthropic.
- Liang, P., et al. (2022). Holistic Evaluation of Language Models (HELM).
- Mazeika, M., et al. (2024). HarmBench: A Standardized Red-Teaming Harness for LLMs.
- Mohamed, S., et al. (2020). Decolonial AI: Decolonial Theory as Sociotechnical Foresight in Artificial Intelligence.
- Pattnayak, P., & Chowdhuri, S. (2026). IndicSafe: A Benchmark for Evaluating Multilingual LLM Safety in South Asia. arXiv:2603.17915.
- Wang, X., et al. (2025). Refusal Direction Is Universal Across Languages. NeurIPS 2025.
- World Economic Forum (2024). Deepfakes: How India Is Tackling Misinformation During Elections.
- Yoo, H., et al. (2025). Code-Switching Red-Teaming: LLM Evaluation for Safety and Multilingual Understanding. ACL 2025.
-

Appendix A: Reproducibility

```
git clone [your-repo-url]
cd safety-hackathon
pip install -r requirements.txt
cp .env.example .env # add OPENAI_API_KEY
python src/build_prompts.py
python src/run_eval.py
python src/analyze.py
```

Results reproduce from `results/eval_results.csv` (288 rows). Analysis outputs in `results/analysis/`.

Appendix B: Figure 1

Figure 1 (`results/analysis/figure1_judge_vs_audited.png`) contrasts the automated LLM-judge ASR (red) against the author-audited strict ASR (blue) across the four registers, for the two India-specific categories. It visualizes both contributions at once: (i) the English-vs-Indic asymmetry — EN is 0% under both label sets; and (ii) the judge over-count — the caste signal (33.3% MONO under the judge) collapses to 0% under audit, while the electoral-misinformation signal largely survives (8.3–16.7%), isolating the one confirmed cross-register bypass. Generated reproducibly from `eval_results.csv` + `harmful_comply_audit.csv` via `python scripts/make_figure.py`.

Figure 1: Judge vs. audited ASR by register — caste signal collapses, electoral-misinfo survives




Figure 1: Automated-judge vs. author-audited ASR by register. Caste signal (33.3% MONO under the judge) collapses to 0% under audit; electoral-misinformation signal largely survives.

Report prepared for submission to the Global South AI Safety Hackathon (Apart Research).
Replace bracketed author fields before PDF export.